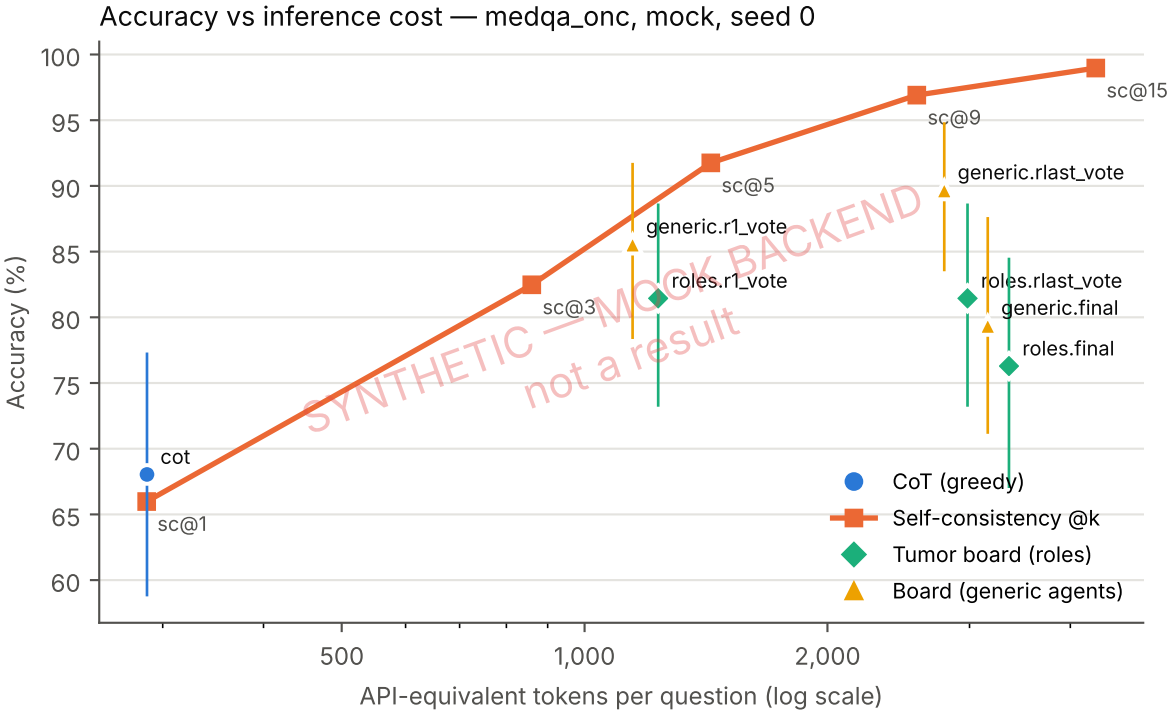
Accuracy vs inference cost — medqa_onc, mock, seed 0
Accuracy (%)
100
95
90
85
80
75
70
65
60
500
1,000
2,000
API-equivalent tokens per question (log scale)
cot
sc@1
sc@3
sc@5
sc@9
sc@15
generic.r1_vote
roles.r1_vote
generic.rlast_vote
roles.rlast_vote
generic.final
roles.final
SYNTHETIC — MOCK BACKEND
not a result
CoT (greedy)
Self-consistency @k
Tumor board (roles)
Board (generic agents)